

PART 3

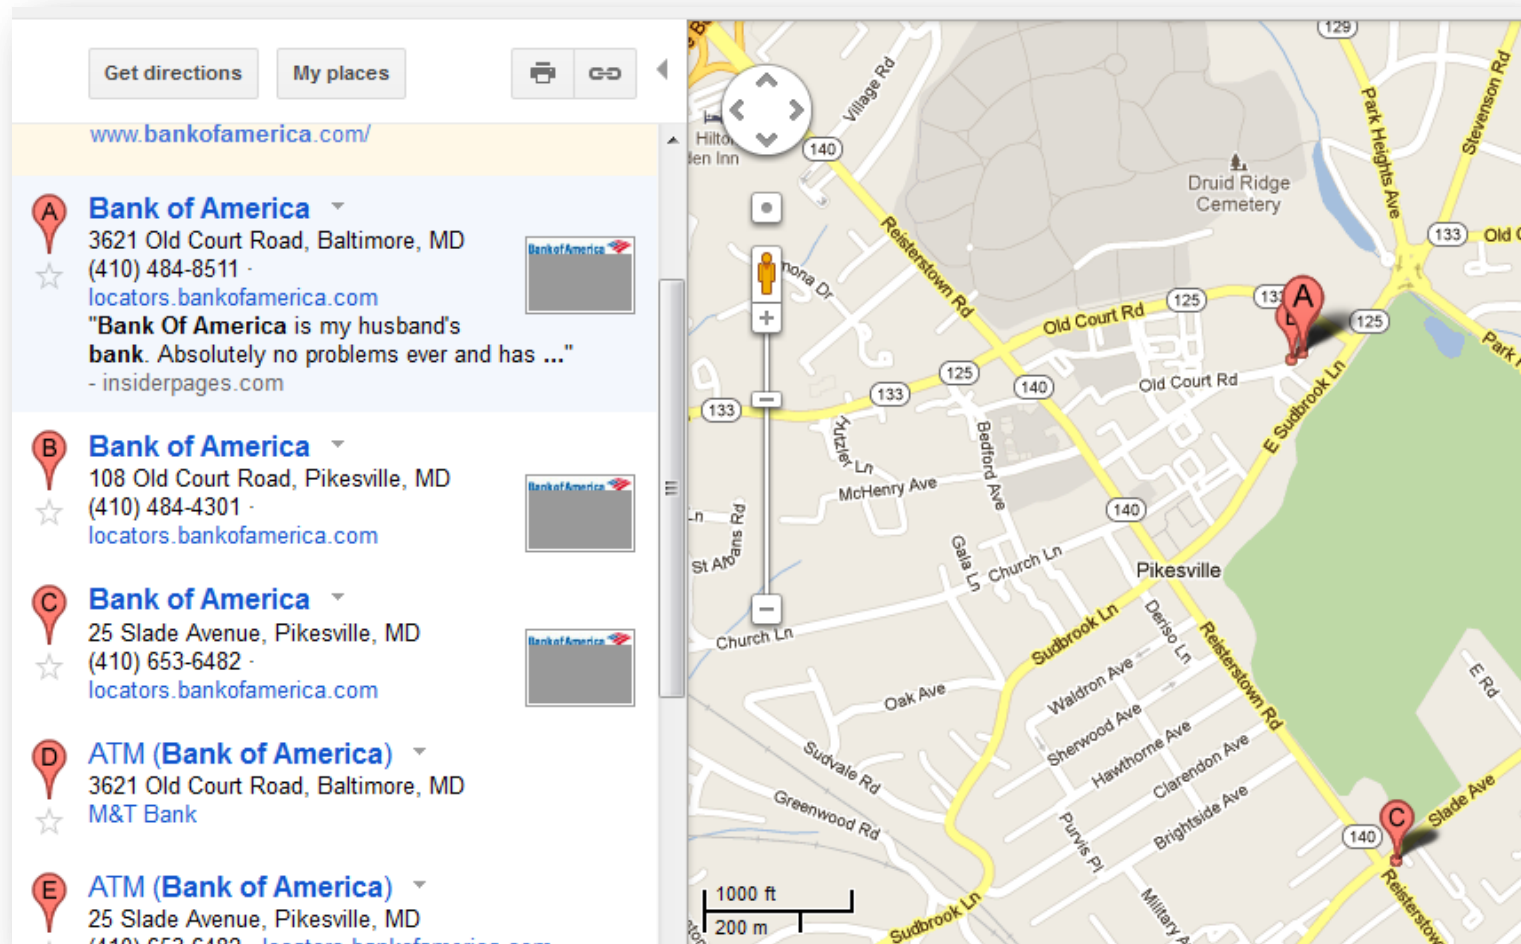
**BLOCKING/CANOPY GENERATION**

# Blocking: Motivation

- Naïve pairwise:  $|R|^2$  pairwise comparisons
  - 1000 business listings each from 1,000 different cities across the world
  - 1 trillion comparisons
  - 11.6 days (if each comparison is 1  $\mu$ s)
- Mentions from different cities are unlikely to be matches
  - 10 million comparisons
  - 10 seconds (if each comparison is 1  $\mu$ s)

# Blocking: Motivation

- Mentions from different cities are unlikely to be matches
  - May miss potential matches



# Blocking: Problem Statement

*Input:* Set of records  $R$

*Output:* Set of *blocks/canopies*

$$\{C_1, C_2, \dots, C_k\}, \text{ where } \forall_i C_i \subset R \text{ and } \bigcup_i C_i = R$$

*Variants:*

- *Disjoint Blocking:* Each mention appears in one block.

$$\forall_{i,j} C_i \cap C_j = \emptyset$$

- *Non-disjoint Blocking:* Mentions can appear in more than one block.

# Blocking: Problem Statement

$\{C_1, C_2, \dots, C_k\}$ , where  $\forall_i C_i \subset R$  and  $\bigcup_i C_i = R$

*Metrics:*

- **Efficiency** (or reduction ratio) : 
$$\frac{\text{number of pairs compared}}{\text{total number of pairs in } R \times R}$$
$$= \frac{|\{(x, y) \mid \exists i C_i, s.t. \ x, y \in C_i\}|}{r(r-1)/2}$$
- **Recall\*** (or pairs completeness) : 
$$\frac{\text{number of true matches compared}}{\text{number of true matches in } R \times R}$$

*\*Need to know ground truth in order to compute this metric*

# Blocking: Problem Statement

## *Metrics:*

- Efficiency (or reduction ratio) :  $\frac{\text{number of pairs compared}}{\text{total number of pairs in } R \times R}$
- Recall\* (or pairs completeness) :  $\frac{\text{number of true matches compared}}{\text{number of true matches in } R \times R}$
- Precision\* (or pairs quality) :  $\frac{\text{number of true matches compared}}{\text{number of matches compared}}$
- Max Canopy Size:  $\max_i |C_i|$

*\*Need to know ground truth in order to compute this metric*

# Blocking Algorithms 1

- Hash based blocking
  - Each block  $C_i$  is associated with a hash key  $h_i$ .
  - Mention  $x$  is hashed to  $C_i$  if  $hash(x) = h_i$ .
  - Within a block, all pairs are compared.
  - Results in disjoint blocks.
- What *hash* function?
  - Deterministic function of attribute values
  - Boolean Functions over attribute values [Bilenko et al ICDM'06, Michelson et al AAAI'06, Das Sarma et al Corr'11]
  - **minHash** (min-wise independent permutations) [Broder et al STOC'98]

# Blocking Algorithms 2

- Pairwise Similarity/Neighborhood based blocking
  - Nearby nodes according to a similarity metric are clustered together
  - Results in non-disjoint canopies.
- Techniques
  - Sorted Neighborhood Approach [Hernandez et al SIGMOD'95]
  - Canopy Clustering [McCallum et al KDD'00]



# Simple Blocking: Inverted Index on a Key

Examples of blocking keys:


- First three characters of last name
- City + State + Zip
- Character or Token n-grams
- Minimum infrequent n-grams

# Learning Optimal Blocking Functions

- Using one or more blocking keys may be insufficient
  - 2,376,206 American's shared the surname Smith in the 2000 US
  - NULL values may create large blocks.
- Solution: Construct blocking functions by combining simple functions


# Complex Blocking Functions

- Conjunction of functions [Michelson et al AAAI'06, Bilenko et al ICDM'06]
  - {City} AND {last four digits of phone}
- Chain-trees [Das Sarma et al Corr '11]
  - **If** ({City} = NULL or LA) **then** {last four digits of phone} AND {area code}  
**else** {last four digits of phone} AND {City}
- BlkTrees [Das Sarma et al Corr '11]



# Learning an Optimal function [Bilenko et al ICDM '06]

- Find  $k$  blocking functions that eliminate the most non-matches, while retaining almost all matches.
  - Need a training set of positive and negative pairs
- Algorithm Idea: Red-Blue Set Cover




Pick  $k$  Blocking keys such that

- (a) At most  $\epsilon$  blue nodes are not covered
- (b) Number of red nodes covered is minimized

# Learning an Optimal function [Bilenko et al ICDM '06]

- Algorithm Idea: Red-Blue Set Cover



Pick  $k$  Blocking keys such that

- (a) At most  $\epsilon$  blue nodes are not covered
- (b) Number of red nodes covered is minimized

- Greedy Algorithm:

- Construct “good” conjunctions of blocking keys  $\{p_1, p_2, \dots\}$ .
- Pick  $k$  conjunctions  $\{p_{i1}, p_{i2}, \dots, p_{ik}\}$ , such that the following is minimized

$$\frac{\text{number of new blue nodes covered by } p_{ij}}{\text{number of red nodes covered by } p_{ij}}$$

# minHash (Minwise Independent Permutations)


- Let  $F_x$  be a set of features for mention  $x$ 
  - (functions of) attribute values
  - character ngrams
  - optimal blocking functions ...
- Let  $\pi$  be a random permutation of features in  $F_x$ 
  - E.g., order imposed by a random hash function
- $\text{minHash}(x)$  = minimum element in  $F_x$  according to  $\pi$

# Why minHash works?

**Surprising property:** For a random permutation  $\pi$ ,


$$P(\text{minHash}(x) = \text{minhash}(y)) = \frac{F_x \cap F_y}{F_x \cup F_y}$$

How to build a blocking scheme such that only pairs with Jacquard similarity  $> s$  fall in the same block (with high prob)?



# Blocking using minHashes

- Compute minHashes using  $r * k$  permutations (hash functions)



- Signature's that match on ***1 out of  $k$***  bands, go to the same block.



# minHash Analysis

False Negatives: (missing matches)

P(pair x,y not in the same block  
with Jacquard sim = s) =  $(1 - s^r)^k$

**should be very low for high similarity pairs**

False Positives: (blocking non-matches)


P(pair x,y in the same block  
with Jacquard sim = s) =  $k \times s^r$

$r = 5, k = 20$

Sim(s)	P(not same block)
0.9	$10^{-8}$
0.8	0.00035
0.7	0.025
0.6	0.2
0.5	0.52
0.4	0.81
0.3	0.95
0.2	0.994
0.1	0.9998

# Sorted Neighborhood [Hernandez et al SIGMOD'95]

- Compute a **Key** for each mention.
- **Sort** the mentions based on the key.
- **Merge**: Check whether a record matches with  $(w-1)$  previous records.
  - Efficient implementation using *Sort Merge Band Join* [DeWitt et al VLDB'91]
- Perform multiple passes with different keys




# Canopy Clustering [McCallum et al KDD'00]

Input: Mentions  $M$ ,  
 $d(x,y)$ , a distance metric,  
thresholds  $T_1 > T_2$

Algorithm:

1. Pick a random element  $x$  from  $M$
2. Create new canopy  $C_x$  using mentions  $y$  s.t.  $d(x,y) < T_1$
3. Delete all mentions  $y$  from  $M$  s.t.  $d(x,y) < T_2$
4. Return to Step 1 if  $M$  is not empty.



# Summary of Blocking

- $O(|R|^2)$  pairwise computations can be prohibitive.
- Blocking eliminates comparisons on a large fraction of non-matches.
- Equality-based Blocking:
  - Construct (one or more) blocking keys from features
  - Records not matching on any key are not compared.
- Similarity based Blocking:
  - Form overlapping canopies of records based on similarity.
  - Only compare records within a cluster.