

Conditional Random Fields: An Introduction*

Hanna M. Wallach

February 24, 2004

1 Labeling Sequential Data

The task of assigning label sequences to a set of observation sequences arises in many fields, including bioinformatics, computational linguistics and speech recognition [6, 9, 12]. For example, consider the natural language processing task of labeling the words in a sentence with their corresponding part-of-speech (POS) tags. In this task, each word is labeled with a tag indicating its appropriate part of speech, resulting in annotated text, such as:

- (1) [PRP He] [VBZ reckons] [DT the] [JJ current] [NN account] [NN deficit] [MD will] [VB narrow] [TO to] [RB only] [# #] [CD 1.8] [CD billion] [IN in] [NNP September] [. .]

Labeling sentences in this way is a useful preprocessing step for higher natural language processing tasks: POS tags augment the information contained within words alone by explicitly indicating some of the structure inherent in language.

One of the most common methods for performing such labeling and segmentation tasks is that of employing hidden Markov models [13] (HMMs) or probabilistic finite-state automata to identify the most likely sequence of labels for the words in any given sentence. HMMs are a form of generative model, that defines a joint probability distribution $p(\mathbf{X}, \mathbf{Y})$ where \mathbf{X} and \mathbf{Y} are random variables respectively ranging over observation sequences and their corresponding label sequences. In order to define a joint distribution of this nature, generative models must enumerate all possible observation sequences – a task which, for most domains, is intractable unless observation elements are represented as isolated units, independent from the other elements in an observation sequence. More precisely, the observation element at any given instant in time may only directly

对一个观察序列标注标签序列的任务比较常见。

在NLP中，标注词性就是这样的。

做这个序列标注的任务方法，主要有HMM和概率有限状态机。

HMM是一个生成模型，它定义了一个联合概率分布 $p(\mathbf{x}, \mathbf{y})$ ，其中 \mathbf{x} ， \mathbf{y} 都是随机变量。

为了得到这样的联合概率分布，生成模型必须枚举所有可能的观察序列（这个任务比较难）。

更进一步说，在任何给定的观察元素中，可能直接依赖状态或者标注。这是一个对一些简单数据点的近似假设。

但是，真实的观察序列通常展示为元素之间有多重相互作用的特征和长期依赖关系。

*University of Pennsylvania CIS Technical Report MS-CIS-04-21

depend on the state, or label, at that time. This is an appropriate assumption for a few simple data sets, however most real-world observation sequences are best represented in terms of multiple interacting features and long-range dependencies between observation elements.

This representation issue is one of the most fundamental problems when labeling sequential data. Clearly, a model that supports tractable inference is necessary, however a model that represents the data without making unwarranted independence assumptions is also desirable. One way of satisfying both these criteria is to use a model that defines a conditional probability $p(\mathbf{Y}|\mathbf{x})$ over label sequences given a particular observation sequence \mathbf{x} , rather than a joint distribution over both label and observation sequences. Conditional models are used to label a novel observation sequence \mathbf{x}_* by selecting the label sequence \mathbf{y}_* that maximizes the conditional probability $p(\mathbf{y}_*|\mathbf{x}_*)$. The conditional nature of such models means that no effort is wasted on modeling the observations, and one is free from having to make unwarranted independence assumptions about these sequences; arbitrary attributes of the observation data may be captured by the model, without the modeler having to worry about how these attributes are related.

Conditional random fields [8] (CRFs) are a probabilistic framework for labeling and segmenting sequential data, based on the conditional approach described in the previous paragraph. A CRF is a form of undirected graphical model that defines a single log-linear distribution over label sequences given a particular observation sequence. The primary advantage of CRFs over hidden Markov models is their conditional nature, resulting in the relaxation of the independence assumptions required by HMMs in order to ensure tractable inference. Additionally, CRFs avoid the label bias problem [8], a weakness exhibited by maximum entropy Markov models [9] (MEMMs) and other conditional Markov models based on directed graphical models. CRFs outperform both MEMMs and HMMs on a number of real-world sequence labeling tasks [8, 11, 15].

2 Undirected Graphical Models

A conditional random field may be viewed as an undirected graphical model, or Markov random field [3], globally conditioned on \mathbf{X} , the random variable representing observation sequences. Formally, we define $G = (V, E)$ to be an undirected graph such that there is a node $v \in V$ corresponding to each of the random variables representing an element Y_v of \mathbf{Y} . If each random variable Y_v obeys the Markov property with respect to G , then (\mathbf{Y}, \mathbf{X}) is a conditional random field. In theory the structure of graph G may be arbitrary, provided it represents the conditional independencies in the label sequences being modeled. However, when modeling sequences, the simplest and most common graph structure encountered is that in which the nodes corresponding to elements of

一个模型能支持良好的推理是必须的，但是如果能够不做不必要的独立性假设，就更好了。

满足这两个标准的一个方法是：在给定一个特殊的观察序列 \mathbf{x} 的时候，定义了条件概率 $p(\mathbf{Y}|\mathbf{x})$ 的模型，而不是一个联合概率分布。

条件模型是用标注序列 \mathbf{y}^* 去标注一个观察序列 \mathbf{x}^* ，使得 $p(\mathbf{y}|\mathbf{x})$ 的概率最大化。

CRF比HMM好的地方就是它不需要条件独立性的假设。

另外，CRF也可以避免标注的 bias 问题，这也是比最大熵马尔科夫模型或者其他条件马尔科夫模型好的原因。

定义一个 $G=(V,E)$ 的无向图，其中存在一个 node 点 v 对应于每一个表示一个 y 的随机变量。

如果每一个随机变量 y 符合马尔科夫性，那么 (Y,X) 就是一个 CRF

\mathbf{Y} form a simple first-order chain, as illustrated in Figure 1.




Figure 1: Graphical structure of a chain-structured CRFs for sequences. The variables corresponding to unshaded nodes are *not* generated by the model.

2.1 Potential Functions

The graphical structure of a conditional random field may be used to factorize the joint distribution over elements Y_v of \mathbf{Y} into a normalized product of strictly positive, real-valued potential functions, derived from the notion of conditional independence.¹ Each potential function operates on a subset of the random variables represented by vertices in G . According to the definition of conditional independence for undirected graphical models, the absence of an edge between two vertices in G implies that the random variables represented by these vertices are conditionally independent given all other random variables in the model. The potential functions must therefore ensure that it is possible to factorize the joint probability such that conditionally independent random variables do not appear in the same potential function. The easiest way to fulfill this requirement is to require each potential function to operate on a set of random variables whose corresponding vertices form a maximal clique within G . This ensures that no potential function refers to any pair of random variables whose vertices are not directly connected and, if two vertices appear together in a clique this relationship is made explicit. In the case of a chain-structured CRF, such as that depicted in Figure 1, each potential function will operate on pairs of adjacent label variables Y_i and Y_{i+1} .

如果图上的两个点之间没有边相连，表示这两个点所代表的随机变量在任何条件下都是条件独立的。

满足这个要求的办法是要求每一个势函数都操作一个随机变量的子集，他们对应的节点是G中的一个最大团。

It is worth noting that an isolated potential function does not have a direct probabilistic interpretation, but instead represents constraints on the configurations of the random variables on which the function is defined. This in turn affects the probability of global configurations – a global configuration with a high probability is likely to have satisfied more of these constraints than a global configuration with a low probability.

¹The product of a set of strictly positive, real-valued functions is not guaranteed to satisfy the axioms of probability. A normalization factor is therefore introduced to ensure that the product of potential functions is a valid probability distribution over the random variables represented by vertices in G .

3 Conditional Random Fields

Lafferty *et al.* [8] define the probability of a particular label sequence \mathbf{y} given observation sequence \mathbf{x} to be a normalized product of potential functions, each of the form

$$\exp\left(\sum_j \lambda_j t_j(y_{i-1}, y_i, \mathbf{x}, i) + \sum_k \mu_k s_k(y_i, \mathbf{x}, i)\right), \quad (2)$$

where $t_j(y_{i-1}, y_i, \mathbf{x}, i)$ is a transition feature function of the entire observation sequence and the labels at positions i and $i-1$ in the label sequence; $s_k(y_i, \mathbf{x}, i)$ is a state feature function of the label at position i and the observation sequence; and λ_j and μ_k are parameters to be estimated from training data.

When defining feature functions, we construct a set of real-valued features $b(\mathbf{x}, i)$ of the observation to express some characteristic of the empirical distribution of the training data that should also hold of the model distribution. An example of such a feature is

$$b(\mathbf{x}, i) = \begin{cases} 1 & \text{if the observation at position } i \text{ is the word "September"} \\ 0 & \text{otherwise.} \end{cases}$$

Each feature function takes on the value of one of these real-valued observation features $b(\mathbf{x}, i)$ if the current state (in the case of a state function) or previous and current states (in the case of a transition function) take on particular values. All feature functions are therefore real-valued. For example, consider the following transition function:

$$t_j(y_{i-1}, y_i, \mathbf{x}, i) = \begin{cases} b(\mathbf{x}, i) & \text{if } y_{i-1} = \text{IN and } y_i = \text{NNP} \\ 0 & \text{otherwise.} \end{cases}$$

In the remainder of this report, notation is simplified by writing

$$s(y_i, \mathbf{x}, i) = s(y_{i-1}, y_i, \mathbf{x}, i)$$

and

$$F_j(\mathbf{y}, \mathbf{x}) = \sum_{i=1}^n f_j(y_{i-1}, y_i, \mathbf{x}, i),$$

where each $f_j(y_{i-1}, y_i, \mathbf{x}, i)$ is either a state function $s(y_{i-1}, y_i, \mathbf{x}, i)$ or a transition function $t(y_{i-1}, y_i, \mathbf{x}, i)$. This allows the probability of a label sequence \mathbf{y} given an observation sequence \mathbf{x} to be written as

$$p(\mathbf{y}|\mathbf{x}, \boldsymbol{\lambda}) = \frac{1}{Z(\mathbf{x})} \exp\left(\sum_j \lambda_j F_j(\mathbf{y}, \mathbf{x})\right). \quad (3)$$

$Z(\mathbf{x})$ is a normalization factor.

4 Maximum Entropy

The form of a CRF, as given in (3), is heavily motivated by the principle of maximum entropy – a framework for estimating probability distributions from a set of training data. Entropy of a probability distribution [16] is a measure of uncertainty and is maximized when the distribution in question is as uniform as possible. The principle of maximum entropy asserts that the only probability distribution that can justifiably be constructed from incomplete information, such as finite training data, is that which has maximum entropy subject to a set of constraints representing the information available. Any other distribution will involve unwarranted assumptions. [7]

If the information encapsulated within training data is represented using a set of feature functions such as those described in the previous section, the maximum entropy model distribution is that which is as uniform as possible while ensuring that the expectation of each feature function with respect to the empirical distribution of the training data equals the expected value of that feature function with respect to the model distribution. Identifying this distribution is a constrained optimization problem that can be shown [2, 10, 14] to be satisfied by (3).

5 Maximum Likelihood Parameter Inference

Assuming the training data $\{(\mathbf{x}^{(k)}, \mathbf{y}^{(k)})\}$ are independently and identically distributed, the product of (3) over all training sequences, as a function of the parameters $\boldsymbol{\lambda}$, is known as the likelihood, denoted by $p(\{\mathbf{y}^{(k)}\}|\{\mathbf{x}^{(k)}\}, \boldsymbol{\lambda})$. Maximum likelihood training chooses parameter values such that the logarithm of the likelihood, known as the log-likelihood, is maximized. For a CRF, the log-likelihood is given by

$$\mathcal{L}(\boldsymbol{\lambda}) = \sum_k \left[\log \frac{1}{Z(\mathbf{x}^{(k)})} + \sum_j \lambda_j F_j(\mathbf{y}^{(k)}, \mathbf{x}^{(k)}) \right].$$

This function is concave, guaranteeing convergence to the global maximum.

Differentiating the log-likelihood with respect to parameter λ_j gives

$$\begin{aligned} \frac{\partial \mathcal{L}(\boldsymbol{\lambda})}{\partial \lambda_j} &= E_{\tilde{p}(\mathbf{Y}, \mathbf{X})} [F_j(\mathbf{Y}, \mathbf{X})] - \\ &\quad \sum_k E_{p(\mathbf{Y}|\mathbf{x}^{(k)}, \boldsymbol{\lambda})} [F_j(\mathbf{Y}, \mathbf{x}^{(k)})], \end{aligned}$$

where $\tilde{p}(\mathbf{Y}, \mathbf{X})$ is the empirical distribution of training data and $E_p[\cdot]$ denotes expectation with respect to distribution p . Note that setting this derivative to

zero yields the maximum entropy model constraint: The expectation of each feature with respect to the model distribution is equal to the expected value under the empirical distribution of the training data.

It is not possible to analytically determine the parameter values that maximize the log-likelihood – setting the gradient to zero and solving for λ does not always yield a closed form solution. Instead, maximum likelihood parameters must be identified using an iterative technique such as iterative scaling [5, 1, 10] or gradient-based methods [15, 17].

6 CRF Probability as Matrix Computations

For a chain-structured CRF in which each label sequence is augmented by start and end states, y_0 and y_{n+1} , with labels **start** and **end** respectively, the probability $p(\mathbf{y}|\mathbf{x}, \lambda)$ of label sequence \mathbf{y} given an observation sequence \mathbf{x} may be efficiently computed using matrices.

Letting \mathcal{Y} be the alphabet from which labels are drawn and y and y' be labels drawn from this alphabet, we define a set of $n + 1$ matrices $\{M_i(\mathbf{x}) | i = 1, \dots, n + 1\}$, where each $M_i(\mathbf{x})$ is a $|\mathcal{Y} \times \mathcal{Y}|$ matrix with elements of the form

$$M_i(y', y|\mathbf{x}) = \exp\left(\sum_j \lambda_j f_j(y', y, \mathbf{x}, i)\right).$$

The unnormalized probability of label sequence \mathbf{y} given observation sequence \mathbf{x} may be written as the product of the appropriate elements of the $n + 1$ matrices for that pair of sequences:

$$p(\mathbf{y}|\mathbf{x}, \lambda) = \frac{1}{Z(\mathbf{x})} \prod_{i=1}^{n+1} M_i(y_{i-1}, y_i|\mathbf{x}).$$

Similarly, the normalization factor $Z(\mathbf{x})$ for observation sequence \mathbf{x} , may be computed from the set of $M_i(\mathbf{x})$ matrices using closed semirings, an algebraic structure that provides a general framework for solving path problems in graphs. Omitting details, $Z(\mathbf{x})$ is given by the (**start**, **end**) entry of the product of all $n + 1$ $M_i(\mathbf{x})$ matrices:

$$Z(\mathbf{x}) = \left[\prod_{i=1}^{n+1} M_i(\mathbf{x}) \right]_{\text{start}, \text{end}} \quad (4)$$

7 Dynamic Programming

In order to identify the maximum-likelihood parameter values – irrespective of whether iterative scaling or gradient-based methods are used – it must be possible to efficiently compute the expectation of each feature function with respect to the CRF model distribution for every observation sequence $\mathbf{x}^{(k)}$ in the training data, given by

$$E_{p(\mathbf{Y}|\mathbf{x}^{(k)}, \boldsymbol{\lambda})} [F_j(\mathbf{Y}, \mathbf{x}^{(k)})] = \sum_{\mathbf{y}} p(\mathbf{Y} = \mathbf{y}|\mathbf{x}^{(k)}, \boldsymbol{\lambda}) F_j(\mathbf{y}, \mathbf{x}^{(k)}). \quad (5)$$

Performing such calculations in a naïve fashion is intractable due to the required sum over label sequences: If observation sequence $\mathbf{x}^{(k)}$ has n elements, there are $n^{|\mathcal{Y}|}$ possible corresponding label sequences. Summing over this number of terms is prohibitively expensive.

Fortunately, the right-hand side of (5) may be rewritten as

$$\sum_{i=1}^n \sum_{y', y} p(Y_{i-1} = y', Y_i = y|\mathbf{x}^{(k)}, \boldsymbol{\lambda}) f_j(y', y, \mathbf{x}^{(k)}), \quad (6)$$

eliminating the need to sum over $n^{|\mathcal{Y}|}$ sequences. Furthermore, a dynamic programming method, similar to the forward-backward algorithm for hidden Markov models, may be used to calculate $p(Y_{i-1} = y', Y_i = y|\mathbf{x}^{(k)}, \boldsymbol{\lambda})$.

Defining forward and backward vectors – $\alpha_i(\mathbf{x})$ and $\beta_i(\mathbf{x})$ respectively – by the base cases

$$\alpha_0(y|\mathbf{x}) = \begin{cases} 1 & \text{if } y = \text{start} \\ 0 & \text{otherwise} \end{cases}$$

and

$$\beta_{n+1}(y|\mathbf{x}) = \begin{cases} 1 & \text{if } y = \text{stop} \\ 0 & \text{otherwise} \end{cases}$$

and the recurrence relations

$$\alpha_i(\mathbf{x})^T = \alpha_{i-1}(\mathbf{x})^T M_i(\mathbf{x})$$

and

$$\beta_i(\mathbf{x}) = M_{i+1}(\mathbf{x}) \beta_{i+1}(\mathbf{x}),$$

the probability of Y_i and Y_{i-1} taking on labels y' and y given observation sequence $\mathbf{x}^{(k)}$ may be written as

$$p(Y_{i-1} = y', Y_i = y|\mathbf{x}^{(k)}, \boldsymbol{\lambda}) = \frac{\alpha_{i-1}(y'|\mathbf{x}) M_i(y', y|\mathbf{x}) \beta_i(y|\mathbf{x})}{Z(\mathbf{x})}.$$

$Z(\mathbf{x})$ is given by the **(start, stop)** entry of the product of all $n+1$ $M_i(\mathbf{x})$ matrices as in (4). Substituting this expression into (6) yields an efficient dynamic programming method for computing feature expectations.

References

- [1] A. L. Berger. The improved iterative scaling algorithm: A gentle introduction, 1997.
- [2] A. L. Berger, S. A. Della Pietra, and V. J. Della Pietra. A maximum entropy approach to natural language processing. *Computational Linguistics*, 22(1):39–71, 1996.
- [3] P. Clifford. Markov random fields in statistics. In Geoffrey Grimmett and Dominic Welsh, editors, *Disorder in Physical Systems: A Volume in Honour of John M. Hammersley*, pages 19–32. Oxford University Press, 1990.
- [4] T. H. Cormen, C. E. Leiserson, and R. L. Rivest. *Introduction to Algorithms*. MIT Press/McGraw-Hill, 1990.
- [5] J. Darroch and D. Ratcliff. Generalized iterative scaling for log-linear models. *The Annals of Mathematical Statistics*, 43:1470–1480, 1972.
- [6] R. Durbin, S. Eddy, A. Krogh, and G. Mitchison. *Biological Sequence Analysis: Probabilistic Models of Proteins and Nucleic Acids*. Cambridge University Press, 1998.
- [7] E. T. Jaynes. Information theory and statistical mechanics. *The Physical Review*, 106(4):620–630, May 1957.
- [8] J. Lafferty, A. McCallum, and F. Pereira. Conditional random fields: probabilistic models for segmenting and labeling sequence data. In *International Conference on Machine Learning*, 2001.
- [9] A. McCallum, D. Freitag, and F. Pereira. Maximum entropy Markov models for information extraction and segmentation. In *International Conference on Machine Learning*, 2000.
- [10] S. Della Pietra, V. Della Pietra, and J. Lafferty. Inducing features of random fields. Technical Report CMU-CS-95-144, Carnegie Mellon University, 1995.
- [11] D. Pinto, A. McCallum, X. Wei, and W. B. Croft. Table extraction using conditional random fields. *Proceedings of the ACM SIGIR*, 2003.
- [12] L. Rabiner and B. H. Juang. *Fundamentals of Speech Recognition*. Prentice Hall Signal Processing Series. Prentice-Hall, Inc., 1993.
- [13] L. R. Rabiner. A tutorial on hidden Markov models and selected applications in speech recognition. *Proceedings of the IEEE*, 77(2):257–285, 1989.
- [14] A. Ratnaparkhi. A simple introduction to maximum entropy models for natural language processing. Technical Report 97-08, Institute for Research in Cognitive Science, University of Pennsylvania, 1997.

- [15] F. Sha and F. Pereira. Shallow parsing with conditional random fields. *Proceedings of Human Language Technology, NAACL 2003*, 2003.
- [16] C. E. Shannon. A mathematical theory of communication. *Bell System Tech. Journal*, 27:379–423 and 623–656, 1948.
- [17] H. M. Wallach. Efficient training of conditional random fields. Master’s thesis, University of Edinburgh, 2002.