

## Data Mining for Imbalanced Datasets: An Overview

Nitesh V. Chawla

Department of Computer Science and Engineering  
University of Notre Dame  
IN 46530, USA  
nchawla@cse.nd.edu

**Summary.** A dataset is imbalanced if the classification categories are not approximately equally represented. Recent years brought increased interest in applying machine learning techniques to difficult “real-world” problems, many of which are characterized by imbalanced data. Additionally the distribution of the testing data may differ from that of the training data, and the true misclassification costs may be unknown at learning time. Predictive accuracy, a popular choice for evaluating performance of a classifier, might not be appropriate when the data is imbalanced and/or the costs of different errors vary markedly. In this Chapter, we discuss some of the sampling techniques used for balancing the datasets, and the performance measures more appropriate for mining imbalanced datasets.

**Key words:** imbalanced datasets, classification, sampling, ROC, cost-sensitive measures, precision and recall

### 45.1 Introduction

The issue with imbalance in the class distribution became more pronounced with the applications of the machine learning algorithms to the real world. These applications range from telecommunications management (Ezawa et al., 1996), bioinformatics (Radivojac et al., 2004), text classification (Lewis and Catlett, 1994, Dumais et al., 1998, Mladeníć and Grobelnik, 1999, Cohen, 1995b), speech recognition (Liu et al., 2004), to detection of oil spills in satellite images (Kubat et al., 1998). The imbalance can be an artifact of class distribution and/or different costs of errors or examples. It has received attention from machine learning and Data Mining community in form of Workshops (Japkowicz, 2000b, Chawla et al., 2003a, Dietterich et al., 2003, Ferri et al., 2004) and Special Issues (Chawla et al., 2004a). The range of papers in these venues exhibited the pervasive and ubiquitous nature of the class imbalance issues faced by the Data Mining community. Sampling methodologies continue to be popular in the research work. However, the research continues to evolve with different applications, as each application provides a compelling problem. One focus of the initial workshops was primarily

the performance evaluation criteria for mining imbalanced datasets. The limitation of the accuracy as the performance measure was quickly established. ROC curves soon emerged as a popular choice (Ferri et al., 2004).

The compelling question, given the different class distributions is: *What is the correct distribution for a learning algorithm?* Weiss and Provost presented a detailed analysis on the effect of class distribution on classifier learning (Weiss and Provost, 2003). Our observations agree with their work that the natural distribution is often not the best distribution for learning a classifier (Chawla, 2003). Also, the imbalance in the data can be more characteristic of “sparseness” in feature space than the class imbalance. Various re-sampling strategies have been used such as random oversampling with replacement, random undersampling, focused oversampling, focused undersampling, oversampling with synthetic generation of new samples based on the known information, and combinations of the above techniques (Chawla et al., 2004b).

In addition to the issue of inter-class distribution, another important problem arising due to the sparsity in data is the distribution of data within each class (Japkowicz, 2001a). This problem was also linked to the issue of small disjuncts in the decision tree learning. Yet another, school of thought is a recognition based approach in the form of a one-class learner. The one-class learners provide an interesting alternative to the traditional discriminative approach, where in the classifier is learned on the target class alone (Japkowicz, 2001b, Juszczak and Duin, 2003, Raskutti and Kowalczyk, 2004, Tax, 2001).

In this chapter<sup>1</sup>, we present a liberal overview of the problem of mining imbalanced datasets with particular focus on performance measures and sampling methodologies. We will present our novel oversampling technique, SMOTE, and its extension in the boosting procedure — SMOTEBoost.

## 45.2 Performance Measure

A classifier is, typically, evaluated by a confusion matrix as illustrated in Figure 45.1 (Chawla et al., 2002). The columns are the Predicted class and the rows are the Actual class. In the confusion matrix,  $TN$  is the number of negative examples correctly classified (True Negatives),  $FP$  is the number of negative examples incorrectly classified as positive (False Positives),  $FN$  is the number of positive examples incorrectly classified as negative (False Negatives) and  $TP$  is the number of positive examples correctly classified (True Positives). Predictive accuracy is defined as  $Accuracy = (TP + TN) / (TP + FP + TN + FN)$ .

However, predictive accuracy might not be appropriate when the data is imbalanced and/or the costs of different errors vary markedly. As an example, consider the classification of pixels in mammogram images as possibly cancerous (Woods et al., 1993). A typical mammography dataset might contain 98% normal pixels and 2% abnormal pixels. A simple default strategy of guessing the majority class would give a predictive accuracy of 98%. The nature of the application requires a fairly high rate of correct detection in the minority class and allows for a small error rate in the majority class in order to achieve this (Chawla et al., 2002). Simple predictive accuracy is clearly not appropriate in such situations.

---


<sup>1</sup> The chapter will utilize excerpts from our published work in various Journals and Conferences. Please see the references for the original publications.

	Predicted Negative	Predicted Positive
Actual Negative	TN	FP
Actual Positive	FN	TP

**Fig. 45.1.** Confusion Matrix

### 45.2.1 ROC Curves

The Receiver Operating Characteristic (ROC) curve is a standard technique for summarizing classifier performance over a range of tradeoffs between true positive and false positive error rates (Swets, 1988). The Area Under the Curve (AUC) is an accepted performance metric for a ROC curve (Bradley, 1997).



**Fig. 45.2.** Illustration of Sweeping out an ROC Curve through under-sampling. Increased under-sampling of the majority (negative) class will move the performance from the lower left point to the upper right.

ROC curves can be thought of as representing the family of best decision boundaries for relative costs of TP and FP. On an ROC curve the X-axis represents  $\%FP = FP / (TN + FP)$  and the Y-axis represents  $\%TP = TP / (TP + FN)$ . The ideal point on the ROC curve would be (0,100), that is all positive examples are classified correctly and no negative examples are misclassified as positive. One way an ROC curve can be swept out is by manipulating the balance of training samples for each class in the training set. Figure 45.2 shows an illustration (Chawla et al., 2002). The line  $y = x$  represents the scenario of randomly guessing the class. A single operating point of a classifier can be chosen from the trade-off between the

%TP and %FP, that is, one can choose the classifier giving the best %TP for an acceptable %FP (Neyman-Pearson method) (Egan, 1975). Area Under the ROC Curve (AUC) is a useful metric for classifier performance as it is independent of the decision criterion selected and prior probabilities. The AUC comparison can establish a dominance relationship between classifiers. If the ROC curves are intersecting, the total AUC is an average comparison between models (Lee, 2000).

The ROC convex hull can also be used as a robust method of identifying potentially optimal classifiers (Provost and Fawcett, 2001). Given a family of ROC curves, the ROC convex hull can include points that are more towards the north-west frontier of the ROC space. If a line passes through a point on the convex hull, then there is no other line with the same slope passing through another point with a larger true positive (TP) intercept. Thus, the classifier at that point is optimal under any distribution assumptions in tandem with that slope (Provost and Fawcett, 2001).

Moreover, distribution/cost sensitive applications can require a ranking or a probabilistic estimate of the instances. For instance, revisiting our mammography data example, a probabilistic estimate or ranking of cancerous cases can be decisive for the practitioner (Chawla, 2003; Maloof, 2003). The cost of further tests can be decreased by thresholding the patients at a particular rank. Secondly, probabilistic estimates can allow one to threshold ranking for class membership at values  $< 0.5$ . The ROC methodology by (Hand, 1997) allows for ranking of examples based on their class memberships — whether a randomly chosen majority class example has a higher majority class membership than a randomly chosen minority class example. It is equivalent to the Wilcoxon test statistic.

#### 45.2.2 Precision and Recall

From the confusion matrix in Figure 45.1, we can derive the expression for *precision* and *recall* (Buckland and Gey, 1994).

$$precision = \frac{TP}{TP + FP}$$

$$recall = \frac{TP}{TP + FN}$$

The main goal for learning from imbalanced datasets is to improve the *recall* without hurting the *precision*. However, *recall* and *precision* goals can be often conflicting, since when increasing the true positive for the minority class, the number of false positives can also be increased; this will reduce the precision. The *F-value* metric is one measure that combines the trade-offs of *precision* and *recall*, and outputs a single number reflecting the “goodness” of a classifier in the presence of rare classes. While ROC curves represent the trade-off between values of TP and FP, the *F-value* represents the trade-off among different values of TP, FP, and FN (Buckland and Gey, 1994). The expression for the *F-value* is as follows:

$$F - value = \frac{(1 + \beta^2) * recall * precision}{\beta^2 * recall + precision}$$

where  $\beta$  corresponds to the relative importance of *precision* vs *recall*. It is usually set to 1.

### 45.2.3 Cost-sensitive Measures

#### Cost Matrix

Cost-sensitive measures usually assume that the costs of making an error are known (Turney, 2000, Domingos, 1999, Elkan, 2001). That is one has a cost-matrix, which defines the costs incurred in false positives and false negatives. Each example,  $x$ , can be associated with a cost  $C(i, j, x)$ , which defines the cost of predicting class  $i$  for  $x$  when the “true” class is  $j$ . The goal is to take a decision to minimize the expected cost. The optimal prediction for  $x$  can be defined as

$$\sum_j P(j|x)C(i, j, x) \quad (45.1)$$

The aforementioned equation requires a computation of conditional probabilities of class  $j$  given feature vector or example  $x$ . While the cost equation is straightforward, we don’t always have a cost attached to making an error. The costs can be different for every example and not only for every type of error. Thus,  $C(i, j)$  is not always  $\equiv$  to  $C(i, j, x)$ .

#### Cost Curves

(Drummond and Holte, 2000) propose cost-curves, where the x-axis represents of the fraction of the positive class in the training set, and the y-axis represents the expected error rate grown on each of the training sets. The training sets for a data set is generated by under (or over) sampling. The error rates for class distributions not represented are construed by interpolation. They define two cost-sensitive components for a machine learning algorithm: 1) producing a variety of classifiers applicable for different distributions and 2) selecting the appropriate classifier for the right distribution. However, when the misclassification costs are known, the x-axis can represent the “probability cost function”, which is the normalized product of  $C(-|+) * P(+)$ ; the y-axis represents the expected cost.

## 45.3 Sampling Strategies

Over and under-sampling methodologies have received significant attention to counter the effect of imbalanced data sets (Solberg and Solberg, 1996, Japkowicz, 2000a, Chawla et al., 2002, Weiss and Provost, 2003, Kubat and Matwin, 1997, Jo and Japkowicz, 2004, Batista et al., 2004, Phua and Alahakoon, 2004, Laurikkala, 2001, Ling and Li, 1998). Various studies in imbalanced datasets have used different variants of over and under sampling, and have presented (sometimes conflicting) viewpoints on usefulness of oversampling versus under-sampling (Chawla, 2003, Maloof, 2003, Drummond and Holte, 2003, Batista et al., 2004).

The random under and over sampling methods have their various short-comings. The random undersampling method can potentially remove certain important examples, and random oversampling can lead to overfitting. However, there has been progression in both the under and over sampling methods. (Kubat and Matwin, 1997) used one-sided selection to selectively undersample the original population. They used Tomek Links (Tomek, 1976) to identify the noisy and borderline examples. They also used the Condensed Nearest Neighbor (CNN) rule (Hart, 1968) to remove examples from the majority class that are far away from the decision border. (Laurikkala, 2001) proposed Neighborhood

Cleaning Rule (NCL) to remove the majority class examples. The author computes three nearest neighbors for each of the ( $E_i$ ) examples in the training set. If  $E_i$  belongs to the majority class, and it is misclassified by its three nearest neighbors, then  $E_i$  is removed. If  $E_i$  belongs to the minority class, and it is misclassified by its three nearest neighbors then the majority class examples among the three nearest neighbors are removed. This approach can reach a computational bottleneck for very large datasets, with a large majority class.

(Japkowicz, 2000a) discussed the effect of imbalance in a dataset. She evaluated three strategies: under-sampling, resampling and a recognition-based induction scheme. She considered two sampling methods for both over and undersampling. Random resampling consisted of oversampling the smaller class at random until it consisted of as many samples as the majority class and “focused resampling” consisted of oversampling only those minority examples that occurred on the boundary between the minority and majority classes. Random under-sampling involved under-sampling the majority class samples at random until their numbers matched the number of minority class samples; focused under-sampling involved under-sampling the majority class samples lying further away. She noted that both the sampling approaches were effective, and she also observed that using the sophisticated sampling techniques did not give any clear advantage in the domain considered. However, her oversampling methodologies did not construct any new examples.

(Ling and Li, 1998) also combined over-sampling of the minority class with under-sampling of the majority class. They used lift analysis instead of accuracy to measure a classifier’s performance. They proposed that the test examples be ranked by a confidence measure and then lift be used as the evaluation criteria. In one experiment, they under-sampled the majority class and noted that the best lift index is obtained when the classes are equally represented. In another experiment, they over-sampled the positive (minority) examples with replacement to match the number of negative (majority) examples to the number of positive examples. The over-sampling and under-sampling combination did not provide significant improvement in the lift index.

We developed a novel oversampling technique called SMOTE (Synthetic Minority Over-sampling TEchnique). It can be essential to provide new related information on the positive class to the learning algorithm, in addition to undersampling the majority class. This was the first attempt to introduce new examples in the training data to enrich the data space and counter the sparsity in the distribution. We will discuss SMOTE in more detail in the subsequent section. We combined SMOTE with undersampling. We used ROC analyses to present the results of our findings.

Batista et al. (Batista et al., 2004) evaluated various sampling methodologies on a variety of datasets with different class distributions. They included various methods in both oversampling and undersampling. They conclude that SMOTE+Tomek and SMOTE+ENN are more applicable and give very good results for datasets with a small number of positive class examples. They also noted that the decision trees constructed from the oversampled datasets are usually very large and complex. This is similar to the observation by (Chawla et al., 2002).

### 45.3.1 Synthetic Minority Oversampling Technique: SMOTE

Over-sampling by replication can lead to similar but more specific regions in the feature space as the decision region for the minority class. This can potentially lead to overfitting on the multiple copies of minority class examples. To overcome the overfitting and broaden the decision region of minority class examples, we introduced a novel technique to generate synthetic examples by operating in “feature space” rather than “data space” (Chawla et al., 2002). The minority class is over-sampled by taking each minority class sample and introducing synthetic

examples along the line segments joining any/all of the  $k$  minority class nearest neighbors. Depending upon the amount of over-sampling required, neighbors from the  $k$  nearest neighbors are randomly chosen. Synthetic samples are generated in the following way: Take the difference between the feature vector (sample) under consideration and its nearest neighbor. Multiply this difference by a random number between 0 and 1, and add it to the feature vector under consideration. This causes the selection of a random point along the line segment between two specific features. This approach effectively forces the decision region of the minority class to become more general. For the nominal cases, we take the majority vote for the nominal value amongst the nearest neighbors. We use the modification of Value Distance Metric (VDM) (Cost and Salzberg, 1993) to compute the nearest neighbors for the nominal valued features.

The synthetic examples cause the classifier to create larger and less specific decision regions, rather than smaller and more specific regions, as typically caused by over-sampling with replication. More general regions are now learned for the minority class rather than being subsumed by the majority class samples around them. The effect is that decision trees generalize better.

SMOTE was tested on a variety of datasets, with varying degrees of imbalance and varying amounts of data in the training set, thus providing a diverse testbed. SMOTE forces focused learning and introduces a bias towards the minority class. On most of the experiments, SMOTE using C4.5 (Quinlan, 1992) and Ripper (Cohen, 1995a) as underlying classifiers outperformed other methods including sampling strategies, Ripper's Loss Ratio, and even Naive Bayes by varying the class priors.

## 45.4 Ensemble-based Methods

Combination of classifiers can be an effective technique for improving prediction accuracy. As one of the most popular combining techniques, boosting (Freund and Schapire, 1996), uses adaptive sampling of instances to generate a highly accurate ensemble of classifiers whose individual global accuracy is only moderate. In boosting, the classifiers in the ensemble are trained serially, with the weights on the training instances adjusted adaptively according to the performance of the previous classifiers. The main idea is that the classification algorithm should concentrate on the instances that are difficult to learn. Boosting has received extensive empirical study (Dietterich, 2000, Bauer and Kohavi, 1999), but most of the published work focuses on improving the accuracy of a weak classifier on datasets with well-balanced class distributions. There has been significant interest in the recent literature for embedding cost-sensitivities in the boosting algorithm. We proposed SMOTEBoost that embeds the SMOTE procedure during boosting iterations. CSB (Ting, 2000) and AdaCost boosting algorithms (Fan et al., 1999) update the weights of examples according to the misclassification costs. On the other side, Rare-Boost (Joshi et al., 2001) updates the weights of the examples differently for all four entries shown in Figure 45.1. Guo and Viktor (Guo and Viktor, 2004) propose another technique that modifies the boosting procedure — DataBoost. As compared to SMOTEBoost, which only focuses on the hard minority class cases, this technique employs a synthetic data generation process for both minority and majority class cases.

In addition to boosting, popular sampling techniques have also been deployed to construct ensembles. Radivojac et al. (Radivojac et al., 2004) combined bagging with oversampling methodologies for the bioinformatics domain. Liu et al. (Liu et al., 2004) also applied a variant of bagging by bootstrapping at equal proportions from both the minority and majority classes. They applied this technique to the problem of sentence boundary detection. Phua et. al (Phua

and Alahakoon, 2004) combine bagging and stacking to identify the best mix of classifiers. In their insurance fraud detection domain, they note that stacking-bagging achieves the best cost-savings

#### 45.4.1 SMOTEBoost

SMOTEBoost algorithm combines SMOTE and the standard boosting procedure (Chawla et al., 2003b). We want to utilize SMOTE for improving the accuracy over the minority classes, and we want to utilize boosting to maintain accuracy over the entire data set. The major goal is to better model the minority class in the data set, by providing the learner not only with the minority class instances that were misclassified in previous boosting iterations, but also with a broader representation of those instances.

The standard boosting procedure gives equal weights to all misclassified examples. Since boosting samples from a pool of data that predominantly consists of the majority class, subsequent samplings of the training set may still be skewed towards the majority class. Although boosting reduces the variance and the bias in the final ensemble (Freund and Schapire, 1996), it might not hold for datasets with skewed class distributions. There is a very strong learning bias towards the majority class cases in a skewed data set, and subsequent iterations of boosting can lead to a broader sampling from the majority class. Boosting (Adaboost) treats both kinds of errors (FP and FN) in a similar fashion. Our goal is to reduce the bias inherent in the learning procedure due to the class imbalance, and increase the sampling weights for the minority class. Introducing SMOTE in each round of boosting will enable each learner to be able to sample more of the minority class cases, and also learn better and broader decision regions for the minority class. SMOTEBoost approach outperformed boosting, Ripper (Cohen, 1995a), and AdaCost on a variety of datasets (Chawla et al., 2003b).

### 45.5 Discussion

Mining from imbalanced datasets is indeed a very important problem from both the algorithmic and performance perspective. Not choosing the right distribution or the objective function while developing a classification model can introduce bias towards majority (potentially uninteresting) class. Furthermore, predictive accuracy is not a useful measure when evaluating classifiers learned on imbalance data sets. Some of the measures discussed in Section 45.2 can be more appropriate.

Sampling methods are very popular in balancing the class distribution before learning a classifier, which uses an error based objective function to search the hypothesis space. We focused on SMOTE in the chapter. Consider the effect on the decision regions in feature space when minority over-sampling is done by replication (sampling with replacement) versus the introduction of synthetic examples. With replication, the decision region that results in a classification decision for the minority class can actually become smaller and more specific as the minority samples in the region are replicated. This is the opposite of the desired effect. Our method of synthetic over-sampling works to cause the classifier to build larger decision regions that contain nearby minority class points. The same reasons may be applicable to why SMOTE performs better than Ripper's loss ratio and Naive Bayes; these methods, nonetheless, are still learning from the information provided in the dataset, albeit with different cost information. SMOTE provides more related minority class samples to learn from, thus allowing a learner to carve broader decision regions, leading to more coverage of the minority class. The



SMOTEBoost methodology that embeds SMOTE within the Adaboost procedure provided further improvements to the minority class prediction.

One compelling problem arising from sampling methodologies is: *Can we identify the right distribution?* Is balanced the best distribution? It is not straightforward. This is very domain and classifier dependent, and is usually driven by empirical observations. (Weiss and Provost, 2003) present a budgeted sampling approach, which represents a heuristic for searching for the right distribution. Another compelling issue is: *What if the test distribution remarkably differs from the training distribution?* If we train a classifier on a distribution tuned on the discovered distribution, will it generalize enough on the testing set. In such cases, one can assume that the natural distribution holds, and apply a form of cost-sensitive learning. If a cost-matrix is known and is static across the training and testing sets, learn from the original or natural distribution, and then apply the cost-matrix at the time of classification. It can also be the case that the majority class is of an equal interest as the minority class — the imbalance here is a mere artifact of class distribution and not of different types of errors (Liu et al., 2004). In such a scenario, it is important to model both the majority and minority classes without a particular bias towards any one class.

We believe mining imbalanced datasets opens a front of interesting problems and research directions. Given that Data Mining is becoming pervasive and ubiquitous in various applications, it is important to investigate along the lines of imbalance both in class distribution and costs.

## Acknowledgements

I would like to thank Larry Hall, Kevin Bowyer and Philip Kegelmeyer for their valuable input during my Ph.D. research in this field. I am also extremely grateful to all my collaborators and co-authors in the area of learning from imbalanced datasets. I have enjoyed working with them and contributing to this field.

## References

- Batista, G. E. A. P. A., Prati, R. C., and Monard, M. C. (2004). A study of the behavior of several methods for balancing machine learning training data. *SIGKDD Explorations*, 6(1).
- Bauer, E. and Kohavi, R. (1999). An empirical comparison of voting classification algorithms: Bagging, boosting and variants. *Machine Learning*, 36(1,2).
- Bradley, A. P. (1997). The Use of the Area Under the ROC Curve in the Evaluation of Machine Learning Algorithms. *Pattern Recognition*, 30(6):1145–1159.
- Buckland, M. and Gey, F. (1994). The Relationship Between Recall and Precision. *Journal of the American Society for Information Science*, 45(1):12–19.
- Chawla, N. V. (2003). C4.5 and Imbalanced Data sets: Investigating the Effect of Sampling Method, Probabilistic Estimate, and Decision Tree Structure. In *ICML Workshop on Learning from Imbalanced Data sets*, Washington, DC.
- Chawla, N. V., Bowyer, K. W., Hall, L. O., and Kegelmeyer, W. P. (2002). SMOTE: Synthetic Minority Oversampling TEchnique. *Journal of Artificial Intelligence Research*, 16:321–357.

- Chawla, N. V., Japkowicz, N., and Kolcz, A., editors (2003a). *Proceedings of the ICML'2003 Workshop on Learning from Imbalanced Data Sets II*.
- Chawla, N. V., Japkowicz, N., and Kolcz, A., editors (2004a). *SIGKDD Special Issue on Learning from Imbalanced Datasets*.
- Chawla, N. V., Japkowicz, N., Kolcz, A. (2004b), Editorial: Learning from Imbalanced Datasets, *SIGKDD Explorations*, 6(1).
- Chawla, N. V., Lazarevic, A., Hall, L. O., and Bowyer, K. W. (2003b). Smoteboost: Improving Prediction of the Minority Class in Boosting. In *Seventh European Conference on Principles and Practice of Knowledge Discovery in Databases*, pages 107–119, Dubrovnik, Croatia.
- Cohen, W. (1995a). Fast Effective Rule Induction. In *Proceedings of the Twelfth International Conference on Machine Learning*, pages 115–123. Department of Computer Science, Katholieke Universiteit Leuven.
- Cohen, W. (1995b). Learning to Classify English Text with ILP Methods. In *Proceedings of the 5th International Workshop on Inductive Logic Programming*, pages 3–24. Department of Computer Science, Katholieke Universiteit Leuven.
- Cost, S. and Salzberg, S. (1993). A Weighted Nearest Neighbor Algorithm for Learning with Symbolic Features. *Machine Learning*, 10(1):57–78.
- Dietterich, T. (2000). An Empirical Comparison of Three Methods for Constructing Ensembles of Decision Trees: Bagging, Boosting and Randomization. *Machine Learning*, 40(2):139 – 157.
- Dietterich, T., Margineantu, D., Provost, F., and Turney, P., editors (2003). *Proceedings of the ICML'2000 Workshop on COST-SENSITIVE LEARNING*.
- Domingos, P. (1999). Metacost: A General Method for Making Classifiers Cost-sensitive. In *Proceedings of the Fifth ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 155–164, San Diego, CA. ACM Press.
- Drummond, C. and Holte, R. (2003). C4.5, class imbalance, and cost sensitivity: Why under-sampling beats over-sampling. In *Proceedings of the ICML'03 Workshop on Learning from Imbalanced Data Sets*.
- Drummond, C. and Holte, R. C. (2000). Explicitly Representing Expected Cost: An Alternative to ROC Representation. In *Proceedings of the Sixth ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 198–207, Boston. ACM.
- Dumais, S., Platt, J., Heckerman, D., and Sahami, M. (1998). Inductive Learning Algorithms and Representations for Text Categorization. In *Proceedings of the Seventh International Conference on Information and Knowledge Management*., pages 148–155.
- Egan, J. P. (1975). Signal Detection Theory and ROC Analysis. In *Series in Cognition and Perception*. Academic Press, New York.
- Elkan, C. (2001). The Foundations of Cost-sensitive Learning. In *Proceedings of the Seventeenth International Joint Conference on Artificial Intelligence*, pages 973–978, Seattle, WA.
- Ezawa, K., J., Singh, M., and Norton, S., W. (1996). Learning Goal Oriented Bayesian Networks for Telecommunications Risk Management. In *Proceedings of the International Conference on Machine Learning, ICML-96*, pages 139–147, Bari, Italy. Morgan Kaufman.
- Fan, W., Stolfo, S., Zhang, J., and Chan, P. (1999). Adacost: Misclassification Cost-sensitive Boosting. In *Proceedings of Sixteenth International Conference on Machine Learning*, pages 983–990, Slovenia.
- Ferri, C., Flach, P., Orallo, J., and Lachice, N., editors (2004). *ECAI' 2004 First Workshop on ROC Analysis in AI*. ECAI.

- Freund, Y. and Schapire, R. (1996). Experiments with a New Boosting Algorithm. In *Thirteenth International Conference on Machine Learning*, Bari, Italy.
- Guo, H. and Viktor, H. L. (2004). Learning from imbalanced data sets with boosting and data generation: The DataBoost-IM approach. *SIGKDD Explorations*, 6(1).
- Hand, D. J. (1997). *Construction and Assessment of Classification Rules*. John Wiley and Sons.
- Hart, P. E. (1968). The Condensed Nearest Neighbor Rule. *IEEE Transactions on Information Theory*, 14:515–516.
- Japkowicz, N. (2000a). The Class Imbalance Problem: Significance and Strategies. In *Proceedings of the 2000 International Conference on Artificial Intelligence (IC-AI'2000): Special Track on Inductive Learning*, Las Vegas, Nevada.
- Japkowicz, N. (2000b). Learning from Imbalanced Data sets: A Comparison of Various Strategies. In *Proceedings of the AAAI'2000 Workshop on Learning from Imbalanced Data Sets*, Austin, TX.
- Japkowicz, N. (2001a). Concept-learning in the presence of between-class and within-class imbalances. In *Proceedings of the Fourteenth Conference of the Canadian Society for Computational Studies of Intelligence*, pages 67–77.
- Japkowicz, N. (2001b). Supervised versus unsupervised binary-learning by feedforward neural networks. *Machine Learning*, 42(1/2):97–122.
- Jo, T. and Japkowicz, N. (2004). Class imbalances versus small disjuncts. *SIGKDD Explorations*, 6(1).
- Joshi, M., Kumar, V., and Agarwal, R. (2001). Evaluating Boosting Algorithms to Classify Rare Classes: Comparison and Improvements. In *Proceedings of the First IEEE International Conference on Data Mining*, pages 257–264, San Jose, CA.
- Juszczak, P. and Duin, R. P. W. (2003). Uncertainty sampling methods for one-class classifiers. In *Proceedings of the ICML'03 Workshop on Learning from Imbalanced Data Sets*.
- Kubat, M., Holte, R., and Matwin, S. (1998). Machine Learning for the Detection of Oil Spills in Satellite Radar Images. *Machine Learning*, 30:195–215.
- Kubat, M. and Matwin, S. (1997). Addressing the Curse of Imbalanced Training Sets: One Sided Selection. In *Proceedings of the Fourteenth International Conference on Machine Learning*, pages 179–186, Nashville, Tennessee. Morgan Kaufmann.
- Laurikkala, J. (2001). Improving Identification of Difficult Small Classes by Balancing Class Distribution. Technical Report A-2001-2, University of Tampere.
- Lee, S. S. (2000). Noisy Replication in Skewed Binary Classification. *Computational Statistics and Data Analysis*, 34.
- Lewis, D. and Catlett, J. (1994). Heterogeneous Uncertainty Sampling for Supervised Learning. In *Proceedings of the Eleventh International Conference of Machine Learning*, pages 148–156, San Francisco, CA. Morgan Kaufmann.
- Ling, C. and Li, C. (1998). Data Mining for Direct Marketing Problems and Solutions. In *Proceedings of the Fourth International Conference on Knowledge Discovery and Data Mining (KDD-98)*, New York, NY. AAAI Press.
- Liu, Y., Chawla, N. V., Shriberg, E., Stolcke, A., and Harper, M. (2004). Resampling Techniques for Sentence Boundary Detection: A Case Study in Machine Learning from Imbalanced Data for Spoken Language Processing. *Under Review*.
- Maloof, M. (2003). Learning when data sets are imbalanced and when costs are unequal and unknown. In *Proceedings of the ICML'03 Workshop on Learning from Imbalanced Data Sets*.

- Mladenić, D. and Grobelnik, M. (1999). Feature Selection for Unbalanced Class Distribution and Naive Bayes. In *Proceedings of the 16th International Conference on Machine Learning*, pages 258–267. Morgan Kaufmann.
- Phua, C. and Alahakoon, D. (2004). Minority report in fraud detection: Classification of skewed data. *SIGKDD Explorations*, 6(1).
- Provost, F. and Fawcett, T. (2001). Robust Classification for Imprecise Environments. *Machine Learning*, 42/3:203–231.
- Quinlan, J. R. (1992). *C4. 5: Programs for Machine Learning*. Morgan Kaufmann, San Mateo, CA.
- Radivojac, P., Chawla, N. V., Dunker, K., and Obradovic, Z. (2004). Classification and Knowledge Discovery in Protein Databases. *Journal of Biomedical Informatics*, 37(4):224–239.
- Raskutti, B. and Kowalczyk, A. (2004). Extreme rebalancing for svms: a case study. *SIGKDD Explorations*, 6(1).
- Solberg, A. H. and Solberg, R. (1996). A Large-Scale Evaluation of Features for Automatic Detection of Oil Spills in ERS SAR Images. In *International Geoscience and Remote Sensing Symposium*, pages 1484–1486, Lincoln, NE.
- Swets, J. (1988). Measuring the Accuracy of Diagnostic Systems. *Science*, 240:1285–1293.
- Tax, D. (2001). *One-class classification*. PhD thesis, Delft University of Technology.
- Ting, K. (2000). A comparative study of cost-sensitive boosting algorithms. In *Proceedings of Seventeenth International Conference on Machine Learning*, pages 983–990, Stanford, CA.
- Tomek, I. (1976). Two Modifications of CNN. *IEEE Transactions on Systems, Man and Cybernetics*, 6:769–772.
- Turney, P. (2000). Types of Cost in Inductive Concept Learning. In *Workshop on Cost-Sensitive Learning at the Seventeenth International Conference on Machine Learning*, pages 15–21, Stanford, CA.
- Weiss, G. and Provost, F. (2003). Learning when Training Data are Costly: The Effect of Class Distribution on Tree Induction. *Journal of Artificial Intelligence Research*, 19:315–354.
- Woods, K., Doss, C., Bowyer, K., Solka, J., Priebe, C., and Kegelmeyer, P. (1993). Comparative Evaluation of Pattern Recognition Techniques for Detection of Microcalcifications in Mammography. *International Journal of Pattern Recognition and Artificial Intelligence*, 7(6):1417–1436.